

An evidence-linked OCR pipeline: text is read from pixels together with its position, page geometry is recovered before language, and specialists are asked only about eligible crops, every component reads and writes one record.




Figure 1: The pipeline runs in one direction. Every stage reads and writes the same positioned evidence record; the two specialist models are invoked only on eligible crops and return alternatives rather than output, and no stage deletes a reading it did not create. Dashed paths fire only when eligibility is met.